

pip install ucimlrepo

Collecting ucimlrepo
 Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.6

import pandas as pd

from ucimlrepo import fetch_ucirepo

fetch dataset
census_income = fetch_ucirepo(id=20)

data (as pandas dataframes)
X = census_income.data.features
y = census_income.data.targets

metadata
print(census_income.metadata)

variable information
print(census_income.variables)

```
{'uci_id': 20, 'name': 'Census Income', 'repository_url': 'https://archive.ics.uci.edu/dataset/20/census+income', 'data_url': 'https://archive.ics.uci.edu/st
name      role      type      demographic \
0      age  Feature  Integer      Age
1      workclass  Feature  Categorical      Income
2      fnlwgt  Feature  Integer      None
3      education  Feature  Categorical  Education Level
4      education-num  Feature  Integer  Education Level
5      marital-status  Feature  Categorical      Other
6      occupation  Feature  Categorical      Other
7      relationship  Feature  Categorical      Other
8      race  Feature  Categorical      Race
9      sex  Feature  Binary      Sex
10     capital-gain  Feature  Integer      None
11     capital-loss  Feature  Integer      None
12     hours-per-week  Feature  Integer      None
13     native-country  Feature  Categorical      Other
14      income  Target  Binary      Income

description units missing_values
0      N/A  None  no
1  Private, Self-emp-not-inc, Self-emp-inc, Feder...  None  yes
2      None  None  no
3  Bachelors, Some-college, 11th, HS-grad, Prof-...  None  no
4      None  None  no
5  Married-civ-spouse, Divorced, Never-married, S...  None  no
6  Tech-support, Craft-repair, Other-service, Sal...  None  yes
7  Wife, Own-child, Husband, Not-in-family, Other...  None  no
8  White, Asian-Pac-Islander, Amer-Indian-Eskimo,...  None  no
9      Female, Male.  None  no
10     None  None  no
11     None  None  no
12     None  None  no
13  United-States, Cambodia, England, Puerto-Rico,...  None  yes
14     >50K, <=50K.  None  no
```

X

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	s
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
...

Next steps: [View recommended plots](#)

y

	income	
0	<=50K	
1	<=50K	
2	<=50K	
3	<=50K	
4	<=50K	
...	...	
48837	<=50K.	
48838	<=50K.	
48839	<=50K.	
48840	<=50K.	
48841	>50K.	
48842 rows × 1 columns		

Next steps: [View recommended plots](#)

```
df = pd.concat((X, y), axis = 1)
df
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
...

Next steps: [View recommended plots](#)

▼ Data cleaning

Checks for null values

```
print(df.isnull().sum())

age          0
workclass    963
fnlwgt       0
education    0
education-num 0
marital-status 0
occupation   966
relationship 0
race         0
sex          0
capital-gain 0
capital-loss 0
hours-per-week 0
native-country 274
income       0
dtype: int64
```

Checks for values on each columns

```
for i in df.columns:
    print(df[i].value_counts(), '\n')
```

```
Japan 92
Guatemala 88
Poland 87
Vietnam 86
Columbia 85
Haiti 75
Portugal 67
Taiwan 65
Iran 59
Greece 49
Nicaragua 49
Peru 46
Ecuador 45
France 38
Ireland 37
Hong 30
Thailand 30
Cambodia 28
Trinidad&Tobago 27
Laos 23
Yugoslavia 23
Outlying-US(Guam-USVI-etc) 23
Scotland 21
Honduras 20
Hungary 19
Holand-Netherlands 1
Name: count, dtype: int64

income
<=50K 24720
<=50K. 12435
>50K 7841
>50K. 3846
Name: count, dtype: int64
```

Saw a '?' values so I wanted to know them

```
for i in df.columns:
    if df[i].dtype == object:
        if (df[i] == '?').any():
            print(i)
            print(df[i][df[i] == '?'].count())
            print()

workclass
1836

occupation
1843

native-country
583
```

Peek on rows with null values

```
null_mask = df.isnull().any(axis=1)
null_rows = df>null_mask

null_rows
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
32565	18	NaN	103497	Some-college	10	Never-married	NaN	Own-child	White	Female
32567	29	NaN	227026	HS-grad	9	Never-married	NaN	Unmarried	Black	Male
32574	58	NaN	299831	HS-grad	9	Married-civ-spouse	NaN	Husband	White	Male
32580	40	Private	85019	Doctorate	16	Married-civ-spouse	Prof-specialty	Husband	Asian-Pac-Islander	Male
32583	72	NaN	132015	7th-8th	4	Divorced	NaN	Not-in-family	White	Female
...
48760	24	NaN	242664	Some-	10	Never-	NaN	Own-child	White	Female

Next steps: [View recommended plots](#)

Checks if when workclass is NaN occupation is also NaN since workclass == 963 and occupation is 966

```
occupation = df[df['occupation'].isnull() & df['workclass'].isnull()]
occupation.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
32565	18	NaN	103497	Some-college	10	Never-married	NaN	Own-child	White	Female
32567	29	NaN	227026	HS-grad	9	Never-married	NaN	Unmarried	Black	Male
						Married-				

Peek on the 3 difference on the workclass and occupation

```
i = df[df['occupation'].isnull() & df['workclass'].notnull()]
i
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
41346	17	Never-worked	131593	11th	7	Never-married	NaN	Own-child	Black	Female

since columns with values of '?' and 'NaN' are the same columns[workclass, occupation, native-country], and We can both conclude or rename them as 'unknown' since they also represent unknown values; to make the representation of data more accurate since it is the data fetched in surveys

```
df.fillna('unknown', inplace=True) # For NaN values
df.replace('?', 'unknown', inplace=True) # For "?" values
```

checks for NaN values and '?' values

```
print(df.isnull().sum())
```

```
age          0
workclass    0
fnlwgt       0
education    0
education-num 0
marital-status 0
occupation   0
relationship 0
race         0
sex          0
capital-gain 0
capital-loss 0
hours-per-week 0
native-country 0
income       0
dtype: int64
```

```
for i in df.columns:
    if df[i].dtype == object:
        if (df[i] == '?').any():
            print(i)
            print(df[i][df[i] == '?'].count())
        else:
            print(f'{i}: None')
```

```
workclass: None
education: None
marital-status: None
occupation: None
relationship: None
race: None
sex: None
native-country: None
income: None
```

Peeks at duplicated values and then drop them

```
dup = df.duplicated().sum()
dup
```

```
29
```

```
df.drop_duplicates(inplace=True)
```

Checks for duplicates

```
dup = df.duplicated().sum()
dup
```

```
0
```

DATA NUMERICAL REPRESENTATION

```
df # peek for the set
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
...

Next steps: [View recommended plots](#)

Checks for dirty uniques

```
for i in df.columns[1:]:
    if df[i].dtype == object:
        print(i)
        print(df[i].unique())
        print("\n")

workclass
['State-gov' 'Self-emp-not-inc' 'Private' 'Federal-gov' 'Local-gov'
 'unknown' 'Self-emp-inc' 'Without-pay' 'Never-worked']

education
['Bachelors' 'HS-grad' '11th' 'Masters' '9th' 'Some-college' 'Assoc-acdm'
 'Assoc-voc' '7th-8th' 'Doctorate' 'Prof-school' '5th-6th' '10th'
 '1st-4th' 'Preschool' '12th']

marital-status
['Never-married' 'Married-civ-spouse' 'Divorced' 'Married-spouse-absent'
 'Separated' 'Married-AF-spouse' 'Widowed']

occupation
['Adm-clerical' 'Exec-managerial' 'Handlers-cleaners' 'Prof-specialty'
 'Other-service' 'Sales' 'Craft-repair' 'Transport-moving'
 'Farming-fishing' 'Machine-op-inspct' 'Tech-support' 'unknown'
 'Protective-serv' 'Armed-Forces' 'Priv-house-serv']

relationship
['Not-in-family' 'Husband' 'Wife' 'Own-child' 'Unmarried' 'Other-relative']

race
['White' 'Black' 'Asian-Pac-Islander' 'Amer-Indian-Eskimo' 'Other']

sex
['Male' 'Female']

native-country
['United-States' 'Cuba' 'Jamaica' 'India' 'unknown' 'Mexico' 'South'
 'Puerto-Rico' 'Honduras' 'England' 'Canada' 'Germany' 'Iran'
 'Philippines' 'Italy' 'Poland' 'Columbia' 'Cambodia' 'Thailand' 'Ecuador'
 'Laos' 'Taiwan' 'Haiti' 'Portugal' 'Dominican-Republic' 'El-Salvador'
 'France' 'Guatemala' 'China' 'Japan' 'Yugoslavia' 'Peru'
 'Outlying-US(Guam-USVI-etc)' 'Scotland' 'Trinidad&Tobago' 'Greece'
 'Nicaragua' 'Vietnam' 'Hong' 'Ireland' 'Hungary' 'Holand-Netherlands']

income
['<=50K' '>50K' '<=50K.' '>50K.']

# Only income is somewhat dirty
print(list(df['income'].unique()))

['<=50K', '>50K', '<=50K.', '>50K.']

df['income'].replace({'<=50K.' : '<=50K', '>50K.' : '>50K'}, inplace = True)

df['income'].value_counts()

income
<=50K      37128
>50K       11685
Name: count, dtype: int64

# Stores the names for later purposes
names = {}
for i in df.columns:
    names[i] = df[i].unique()
# print(names['income'])

print(df['workclass'].unique())
print(names['workclass'])
```

```
['State-gov' 'Self-emp-not-inc' 'Private' 'Federal-gov' 'Local-gov'
 'unknown' 'Self-emp-inc' 'Without-pay' 'Never-worked']
['State-gov' 'Self-emp-not-inc' 'Private' 'Federal-gov' 'Local-gov'
 'unknown' 'Self-emp-inc' 'Without-pay' 'Never-worked']
```

```
df[['education', 'education-num']].value_counts()
```

education	education-num	
HS-grad	9	15777
Some-college	10	10869
Bachelors	13	8020
Masters	14	2656
Assoc-voc	11	2060
11th	7	1812
Assoc-acdm	12	1601
10th	6	1389
7th-8th	4	954
Prof-school	15	834
9th	5	756
12th	8	656
Doctorate	16	594
5th-6th	3	508
1st-4th	2	245
Preschool	1	82

Name: count, dtype: int64

```
df['workclass'].value_counts()
```

workclass	
Private	33879
Self-emp-not-inc	3861
Local-gov	3136
unknown	2799
State-gov	1981
Self-emp-inc	1694
Federal-gov	1432
Without-pay	21
Never-worked	10

Name: count, dtype: int64

```
df['income'].value_counts()
```

income	
<=50K	37128
>50K	11685

Name: count, dtype: int64

```
df['income'].value_counts()
```

income	
<=50K	37128
>50K	11685

Name: count, dtype: int64

```
print(df['native-country'].unique())
print(names['native-country'])
```

```
['United-States' 'Cuba' 'Jamaica' 'India' 'unknown' 'Mexico' 'South'
 'Puerto-Rico' 'Honduras' 'England' 'Canada' 'Germany' 'Iran'
 'Philippines' 'Italy' 'Poland' 'Columbia' 'Cambodia' 'Thailand' 'Ecuador'
 'Laos' 'Taiwan' 'Haiti' 'Portugal' 'Dominican-Republic' 'El-Salvador'
 'France' 'Guatemala' 'China' 'Japan' 'Yugoslavia' 'Peru'
 'Outlying-US(Guam-USVI-etc)' 'Scotland' 'Trinidad&Tobago' 'Greece'
 'Nicaragua' 'Vietnam' 'Hong' 'Ireland' 'Hungary' 'Holand-Netherlands']
['United-States' 'Cuba' 'Jamaica' 'India' 'unknown' 'Mexico' 'South'
 'Puerto-Rico' 'Honduras' 'England' 'Canada' 'Germany' 'Iran'
 'Philippines' 'Italy' 'Poland' 'Columbia' 'Cambodia' 'Thailand' 'Ecuador'
 'Laos' 'Taiwan' 'Haiti' 'Portugal' 'Dominican-Republic' 'El-Salvador'
 'France' 'Guatemala' 'China' 'Japan' 'Yugoslavia' 'Peru'
 'Outlying-US(Guam-USVI-etc)' 'Scotland' 'Trinidad&Tobago' 'Greece'
 'Nicaragua' 'Vietnam' 'Hong' 'Ireland' 'Hungary' 'Holand-Netherlands']
```

```
df['native-country'].value_counts()
```

native-country	
United-States	43810
Mexico	947
unknown	856
Philippines	295
Germany	206
Puerto-Rico	184
Canada	182
El-Salvador	155
India	151
Cuba	138
England	127
China	122
South	115
Jamaica	106
Italy	105
Dominican-Republic	103
Japan	92
Poland	87
Guatemala	86
Vietnam	86
Columbia	85
Haiti	75
Portugal	67
Taiwan	65
Iran	59
Greece	49
Nicaragua	49
Peru	46
Ecuador	45
France	38
Ireland	37
Hong	30
Thailand	30
Cambodia	28
Trinidad&Tobago	27
Laos	23

Yugoslavia	23
Outlying-US(Guam-USVI-etc)	23
Scotland	21
Honduras	20
Hungary	19
Holand-Netherlands	1
Name: count, dtype: int64	

```
for i in df.columns:
    if df[i].dtype == object:
        if i != 'education' and i != 'workclass':
            num = 1
            for j in df[i].unique():
                df[i].replace({j:num}, inplace = True)
            num += 1
```

```
df['education'].replace({'Preschool': 'Elem-grad', '1st-4th': 'Elem-grad', '5th-6th': 'Elem-grad',
                        '12th': 'HS-grad', '9th': 'Elem-grad', '7th-8th': 'Elem-grad',
                        '10th': 'Elem-grad', '11th': 'HS-grad'}, inplace=True)
```

```
df[['education', 'education-num']].value_counts()
```

education	education-num	
HS-grad	9	15777
Some-college	10	10869
Bachelors	13	8020
Masters	14	2656
Assoc-voc	11	2060
11th	7	1812
Assoc-acdm	12	1601
10th	6	1389
7th-8th	4	954
Prof-school	15	834
9th	5	756
12th	8	656
Doctorate	16	594
5th-6th	3	508
1st-4th	2	245
Preschool	1	82
Name: count, dtype: int64		

```
df['native-country'].value_counts() # Chekcs if it is the same with the previous counting and yes it is
```



▼ Data Analysis

```
income_counts = df['income'].value_counts()
percentage_1 = (income_counts[1] / len(df)) * 100
percentage_2 = (income_counts[2] / len(df)) * 100

print("Percentage of 1:", round(percentage_1, 1))
print("Percentage of 2:", round(percentage_2, 1))
```

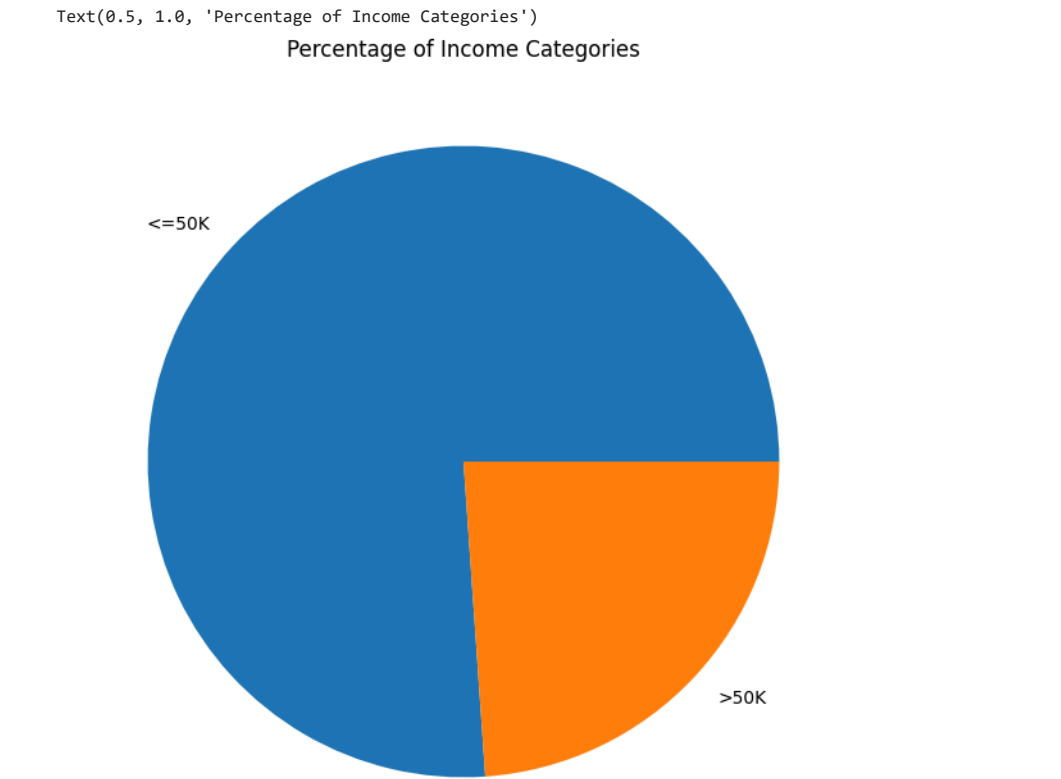
Percentage of 1: 76.1
Percentage of 2: 23.9

```
percent = pd.DataFrame({'Income' : names['income'], 'Count' : df['income'].value_counts().values, 'Percentage' : [round(percentage_1, 1), round(percentage_2, 1)]})
percent
```


	Income	Count	Percentage	
0	<=50K	37128	76.1	
1	>50K	11685	23.9	

Next steps: [View recommended plots](#)

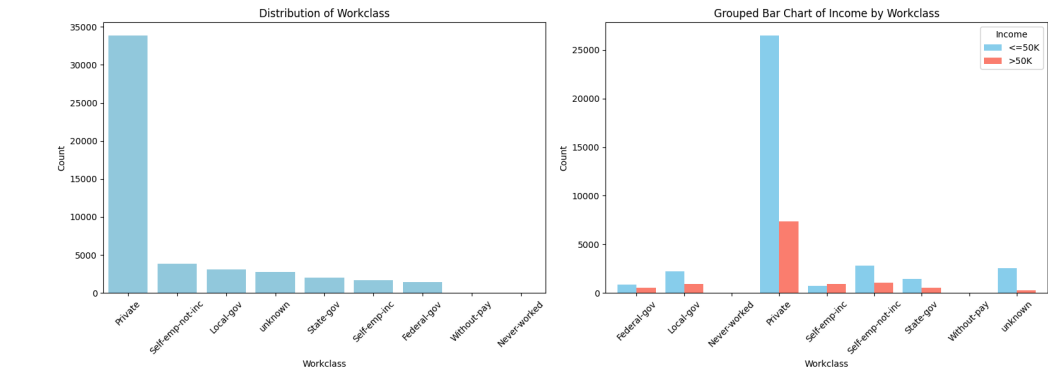
```
percent = [round(percentage_1, 1), round(percentage_2, 1)]
plt.figure(figsize=(8, 8))
plt.pie(percent, labels=names['income'])
plt.title('Percentage of Income Categories')
```



```
import numpy as np
age_groups = np.arange(10, 101, 10)
labels = [f'{age}-{age+9}' for age in age_groups[:-1]]
age_bins = pd.cut(df['age'], bins=age_groups, labels=labels, right=False)
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
sns.barplot(x=age_bins.value_counts().index, y=age_bins.value_counts().values, color='skyblue', ax=axes[0])
axes[0].set_title('Distribution of Age Groups')
axes[0].set_ylabel('Count')
axes[0].set_xlabel('Age Group')
axes[0].tick_params(axis='x', rotation=45)
grouped_counts = df.groupby([age_bins, 'income']).size().unstack()
grouped_counts.plot(kind='bar', stacked=True, color=['skyblue', 'salmon'], ax=axes[1])
axes[1].set_title('Grouped Bar Chart of Income by Age')
axes[1].set_xlabel('Age Group')
axes[1].set_ylabel('Count')
axes[1].legend(labels=names['income'], title='Income')
axes[1].tick_params(axis='x', rotation=45)
plt.tight_layout()
plt.show()
```

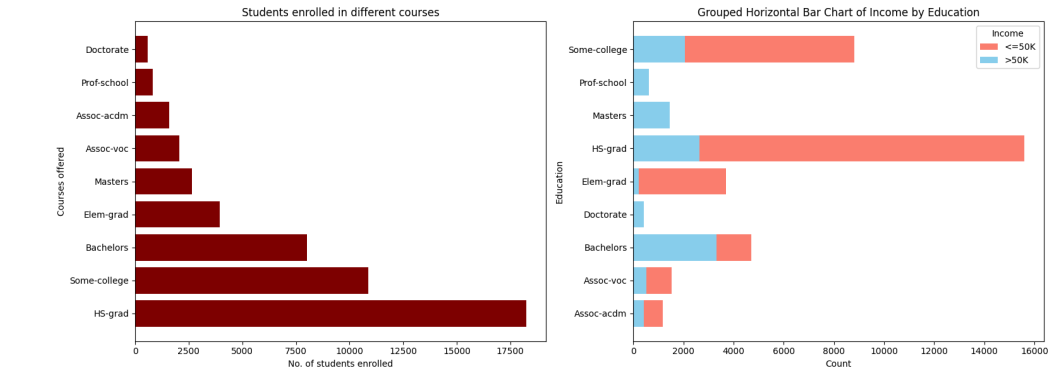



```
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
sns.barplot(x=df['workclass'].value_counts().index, y=df['workclass'].value_counts(), color='skyblue', ax=axes[0])
axes[0].set_title('Distribution of Workclass')
axes[0].set_ylabel('Count')
axes[0].set_xlabel('Workclass')
axes[0].tick_params(axis='x', rotation=45)
grouped_counts = df.groupby(['workclass', 'income']).size().unstack()
grouped_counts.plot(kind='bar', stacked=False, width=0.8, color=['skyblue', 'salmon'], ax=axes[1])
axes[1].legend(labels=names['income'], title='Income')
axes[1].set_title('Grouped Bar Chart of Income by Workclass')
axes[1].set_xlabel('Workclass')
axes[1].set_ylabel('Count')
axes[1].tick_params(axis='x', rotation=45)
plt.tight_layout()
plt.show()
```

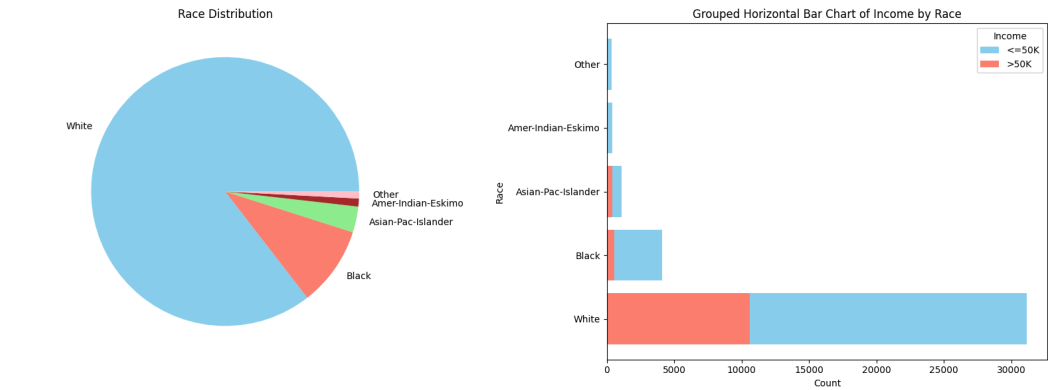


```
import matplotlib.pyplot as plt
```


```
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
course_counts = df['education'].value_counts()
axes[0].barh(course_counts.index, course_counts.values, color='maroon')
axes[0].set_xlabel("No. of students enrolled")
axes[0].set_ylabel("Courses offered")
axes[0].set_title("Students enrolled in different courses")
grouped_counts = df.groupby(['education', 'income']).size().unstack()
education_levels = grouped_counts.index
colors = ['salmon', 'skyblue']
for income_category, color in zip(grouped_counts.columns, colors):
    axes[1].barh(education_levels, grouped_counts[income_category], color=color, label=income_category)
axes[1].legend(labels=names['income'], title='Income')
axes[1].set_title('Grouped Horizontal Bar Chart of Income by Education')
axes[1].set_xlabel('Count')
axes[1].set_ylabel('Education')
plt.tight_layout()
plt.show()
```



```
custom_colors = ['skyblue', 'salmon', 'lightgreen', 'brown', 'pink']
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
axes[0].pie(df['race'].value_counts(), labels=names['race'], colors=custom_colors)
axes[0].set_title('Race Distribution')
grouped_counts = df.groupby(['race', 'income']).size().unstack()
race_names = names['race']
for income_category, color in zip(grouped_counts.columns, custom_colors[:2]):
    axes[1].barh(race_names, grouped_counts[income_category], color=color, label=income_category)
axes[1].set_title('Grouped Horizontal Bar Chart of Income by Race')
axes[1].set_xlabel('Count')
axes[1].set_ylabel('Race')
axes[1].legend(labels=names['income'], title='Income')
plt.tight_layout()
plt.show()
```



```
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
axes[0].pie(df['occupation'].value_counts(), labels=names['occupation'])
axes[0].set_title('Occupation Distribution')
grouped_counts = df.groupby(['occupation', 'income']).size().unstack()
for income_category, color in zip(grouped_counts.columns, ['salmon', 'skyblue']):
    axes[1].barh(names['occupation'], grouped_counts[income_category], color=color, label=income_category)
axes[1].legend(labels=names['income'], title='Income')
axes[1].set_title('Grouped Horizontal Bar Chart of Income by Occupation')
axes[1].set_xlabel('Count')
axes[1].set_ylabel('Occupation')
plt.tight_layout()
plt.show()
```



Sex Distribution

A pie chart titled 'Sex Distribution' showing the proportion of males and females. The male segment is blue and represents approximately 75% of the total, while the female segment is orange and represents approximately 25%.

Sex	Proportion
Male	~75%
Female	~25%

Grouped Horizontal Bar Chart of Income by Sex

A grouped horizontal bar chart titled 'Grouped Horizontal Bar Chart of Income by Sex'. The y-axis is labeled 'Sex' with categories 'Female' and 'Male'. The x-axis represents income. A legend indicates two income groups: '<=50K' (red) and '>50K' (blue). For females, the blue segment is small, while the red segment is large. For males, the blue segment is significantly larger than the red segment.

Sex	<=50K	>50K
Female	~75%	~25%
Male	~45%	~55%

The left chart, titled "Distribution of Hours per Week", is a pie chart showing the distribution of weekly hours. The largest slice is blue, labeled "40". Other significant slices are orange (labeled "13"), green (labeled "16"), red (labeled "45"), purple (labeled "50"), brown (labeled "80"), pink (labeled "30"), grey (labeled "35"), cyan (labeled "60"), and dark blue (labeled "20"). A large number of very thin slices are visible, representing less frequent hour counts.

The right chart, titled "Income Distribution by Hours per Week", is a horizontal bar chart. The y-axis is labeled "Hours per Week" and ranges from 0 to 100. The x-axis is labeled "Count" and ranges from 0 to 12500. The chart shows the count of people for each hour value. The bar for 40 hours is the longest, extending to a count of approximately 12500. Other bars are visible for 20, 30, 50, 60, and 80 hours.

[illegible]